

Assignment3

July 10, 2025

```
[1]: # Importing Libraries
import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import KFold, cross_val_score, train_test_split
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler
from sklearn import metrics
import seaborn as sns

[2]: # Model Training Function
def model_training(X_train, y_train, X_test, y_test, cv=None, penalty=None, C=0.
    ↪1, solver='lbfgs'):
    classifier = LogisticRegression(penalty=penalty, C=C, solver=solver, ↪
    ↪random_state=0)
    classifier.fit(X_train, y_train)
    y_pred = classifier.predict(X_test)
    cnf_matrix = confusion_matrix(y_test, y_pred)
    return cnf_matrix, y_pred

[3]: # Confusion Matrix Plotting Function
def get_confusion_matrix(cnf_matrix):
    class_names = [0, 1] # number of classes
    fig, ax = plt.subplots()
    tick_marks = np.arange(len(class_names))
    plt.xticks(tick_marks, class_names)
    plt.yticks(tick_marks, class_names)

    # create heatmap
    sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu", fmt='g')
    ax.xaxis.set_label_position("top")
    plt.tight_layout()
    plt.title('Confusion matrix', y=1.1)
```

```
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

0.1 Diabetes Dataset

```
[5]: # Read and Display the Diabetes Dataset
df = pd.read_csv("/content/diabetes.csv")
print("CSV File Shape")
print(df.shape)
df.head()
```

CSV File Shape
(768, 9)

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[6]: # Splitting Features and Labels
X = df.iloc[:, :-1].values # All columns except the last one (features)
Y = df.iloc[:, -1].values # Last column (target)
```

```
[7]: # Standardizing the Features
sc_X = StandardScaler()
X_scaled = sc_X.fit_transform(X)
```

```
[8]: # Splitting into Training and Testing Sets
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.20,
↳ random_state=42)
```

Problem 1:


```
[9]: # Train & Evaluate Logistic Regression Model
cnf_matrix, Y_pred = model_training(X_train, y_train, X_test, y_test,
↳ penalty=None)
```

```
[10]: def get_results(y_test, y_pred):
    acc = metrics.accuracy_score(y_test, y_pred)
    print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
    print("Precision:", metrics.precision_score(y_test, y_pred))
    print("Recall:", metrics.recall_score(y_test, y_pred))
    print("F1-score:", metrics.f1_score(y_test, y_pred))
    return acc * 100

# Print Evaluation Metrics
acc_1 = get_results(y_test, Y_pred)
```

Accuracy: 0.7727272727272727
Precision: 0.6851851851851852
Recall: 0.6727272727272727
F1-score: 0.6788990825688074

```
[11]: # Display Confusion Matrix
get_confusion_matrix(cnf_matrix)
```



Cancer Dataset

```
[12]: import pandas as pd

# Load CSV file
df = pd.read_csv("/content/cancer.csv")
print("CSV Shape:", df.shape)
df.head()
```

CSV Shape: (569, 33)

```
[12]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.3001	0.14710	
1	0.08474	0.07864	0.0869	0.07017	
2	0.10960	0.15990	0.1974	0.12790	
3	0.14250	0.28390	0.2414	0.10520	
4	0.10030	0.13280	0.1980	0.10430	

...	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	...	17.33	184.60	2019.0	0.1622
1	...	23.41	158.80	1956.0	0.1238
2	...	25.53	152.50	1709.0	0.1444
3	...	26.50	98.87	567.7	0.2098
4	...	16.67	152.20	1575.0	0.1374

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	\
0	0.6656	0.7119	0.2654	0.4601	
1	0.1866	0.2416	0.1860	0.2750	
2	0.4245	0.4504	0.2430	0.3613	
3	0.8663	0.6869	0.2575	0.6638	
4	0.2050	0.4000	0.1625	0.2364	

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN
2	0.08758	NaN
3	0.17300	NaN
4	0.07678	NaN

[5 rows x 33 columns]

```
[13]: # Drop non-numeric 'id' and 'diagnosis' from features
X = df.drop(['id', 'diagnosis'], axis=1).values

# Encode 'diagnosis' (M = 1, B = 0)
Y = df['diagnosis'].map({'M': 1, 'B': 0}).values
```

```
[14]: # Impute missing values
from sklearn.impute import SimpleImputer
imputer = SimpleImputer(strategy='mean')
X = imputer.fit_transform(X)
```

```
[15]: # Split and Scale the Data
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.20,
    ↪random_state=42)

sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)
```

Problem 2.1

```
[16]: # Train logistic regression
from sklearn.linear_model import LogisticRegression

model = LogisticRegression(penalty='l2', solver='lbfgs', max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```


```
[19]: # Evaluate the Model
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↪f1_score

# Print metrics
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("Accuracy:", acc)
print("Precision:", prec)
print("Recall:", rec)
print("F1-Score:", f1)
```

Accuracy: 0.9736842105263158
Precision: 0.9761904761904762
Recall: 0.9534883720930233
F1-Score: 0.9647058823529412

```
[18]: # Plot Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(pd.DataFrame(cm), annot=True, fmt='g', cmap='YlGnBu')
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("Actual Label")
plt.show()
```




Problem 2.2

```
[20]: # Add the weight penalty and repeat the training
cnf_matrix_2_b, y_pred_2_b = model_training(
    X_train, y_train, X_test, y_test,
    penalty='l2', C=10, solver='lbfgs'
)
```

```
[21]: # Plot Confusion Matrix
acc_2_b = get_results(y_test, y_pred_2_b)
get_confusion_matrix(cnf_matrix_2_b)
```

Accuracy: 0.9736842105263158
Precision: 0.9545454545454546
Recall: 0.9767441860465116
F1-score: 0.9655172413793104



Problem 3:

```
[22]: # Model Training
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix

def model_training_NB(X_train, y_train, X_test, y_test):
    classifier = GaussianNB()
```

```

y_pred = classifier.fit(X_train, y_train).predict(X_test)
cnf_matrix = confusion_matrix(y_test, y_pred)
return cnf_matrix, y_pred

```

```

[23]: # Train and evaluate the model
cnf_matrix_3, y_pred_3 = model_training_NB(X_train, y_train, X_test, y_test)
acc_3 = get_results(y_test, y_pred_3)
get_confusion_matrix(cnf_matrix_3)


```

Accuracy: 0.9649122807017544

Precision: 0.975609756097561

Recall: 0.9302325581395349

F1-score: 0.9523809523809523



Problem 4 & 5:


```
[24]: # PCA + Classification
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

def get_results(y_test, y_pred):
    acc = accuracy_score(y_test, y_pred)
    pre = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    fscore = f1_score(y_test, y_pred)
    print("Accuracy:", acc)
    print("Precision:", pre)
    print("Recall:", rec)
    print("F1-Score:", fscore)
    return [acc*100.0, pre*100.0, rec*100.0, fscore*100.0]

[25]: # Logistic Regression with PCA
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

def logist_model_training_pca(X, Y):
    n = X.shape[1]
    acc_list = []
    recall_list = []
    precision_list = []
    f1score_list = []
    k_list = []

    for i in range(n):
        print("K = " + str(i+1))
        pca = PCA(n_components=i+1)
        principalComponents = pca.fit_transform(X)
        X_train, X_test, y_train, y_test = train_test_split(principalComponents, Y, test_size=0.20, random_state=9999)

        classifier = LogisticRegression(random_state=9)
        y_pred = classifier.fit(X_train, y_train).predict(X_test)

        re = get_results(y_test, y_pred)
        acc_list.append(re[0])
        precision_list.append(re[1])
        recall_list.append(re[2])
        f1score_list.append(re[3])
        k_list.append(i+1)

    high_acc = max(acc_list)
    high_acc_k = acc_list.index(high_acc) + 1
```

```

print("-----")
print("Highest Classification Accuracy Achieved: " + str(high_acc) + " for_
↪K number = " + str(high_acc_k))
return k_list, acc_list, precision_list, recall_list, f1score_list

```

```

[26]: # Naive Bayes with PCA
from sklearn.naive_bayes import GaussianNB

def GaussianNB_model_training_pca(X, Y):
    n = X.shape[1]
    acc_list = []
    recall_list = []
    precision_list = []
    f1score_list = []
    k_list = []

    for i in range(n):
        print("K = " + str(i+1))
        pca = PCA(n_components=i+1)
        principalComponents = pca.fit_transform(X)
        X_train, X_test, y_train, y_test =_
↪train_test_split(principalComponents, Y, test_size=0.20, random_state=9999)

        classifier = GaussianNB()
        y_pred = classifier.fit(X_train, y_train).predict(X_test)

        re = get_results(y_test, y_pred)
        acc_list.append(re[0])
        precision_list.append(re[1])
        recall_list.append(re[2])
        f1score_list.append(re[3])
        k_list.append(i+1)

    high_acc = max(acc_list)
    high_acc_k = acc_list.index(high_acc) + 1
    print("-----")
    print("Highest Classification Accuracy Achieved: " + str(high_acc) + " for_
↪K number = " + str(high_acc_k))
    return k_list, acc_list, precision_list, recall_list, f1score_list

```

```

[27]: # Plotting the Results
import matplotlib.pyplot as plt

def plot_result_with_k(k_list, acc_list, precision_list, recall_list,
↪f1score_list):
    plt.plot(k_list, acc_list, label = "Accuracy")
    plt.plot(k_list, precision_list, label = "Precision")

```

```

plt.plot(k_list, recall_list, label = "Recall")
plt.plot(k_list, f1score_list, label = "F1-Score")
plt.legend()
plt.title("Plotting classification accuracy, precision, recall and F1-score_
↳over a different number of Ks")
plt.xlabel("K")
plt.ylabel("Value")
plt.show()

```

```

[28]: # Loading and Preprocessing Cancer Dataset
import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Load CSV file
df = pd.read_csv("/content/cancer.csv")
print("CSV Shape:", df.shape)
df.head()

# Drop ID and diagnosis column, and extract labels
X = df.drop(['id', 'diagnosis'], axis=1).values
Y = df['diagnosis'].map({'M': 1, 'B': 0}).values

# Impute missing values
imputer = SimpleImputer(strategy='mean')
X = imputer.fit_transform(X)

# Split and scale
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.20,
↳random_state=42)

sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)

```

CSV Shape: (569, 33)

Problem 4:

```

[29]: # perform PCA + logistic regression
k_list, acc_list, precision_list, recall_list, f1score_list =
↳logist_model_training_pca(X, Y)

```

K = 1
Accuracy: 0.9122807017543859
Precision: 0.90625
Recall: 0.8055555555555556

F1-Score: 0.8529411764705882
K = 2
Accuracy: 0.9473684210526315
Precision: 0.9166666666666666
Recall: 0.9166666666666666
F1-Score: 0.9166666666666666
K = 3
Accuracy: 0.956140350877193
Precision: 0.9428571428571428
Recall: 0.9166666666666666
F1-Score: 0.9295774647887324
K = 4
Accuracy: 0.9473684210526315
Precision: 0.8947368421052632
Recall: 0.9444444444444444
F1-Score: 0.918918918918919
K = 5
Accuracy: 0.9473684210526315
Precision: 0.8947368421052632
Recall: 0.9444444444444444
F1-Score: 0.918918918918919
K = 6
Accuracy: 0.9473684210526315
Precision: 0.8947368421052632
Recall: 0.9444444444444444
F1-Score: 0.918918918918919
K = 7
Accuracy: 0.9473684210526315
Precision: 0.8947368421052632
Recall: 0.9444444444444444
F1-Score: 0.918918918918919
K = 8
Accuracy: 0.9473684210526315
Precision: 0.8947368421052632
Recall: 0.9444444444444444
F1-Score: 0.918918918918919
K = 9
Accuracy: 0.9473684210526315
Precision: 0.8947368421052632
Recall: 0.9444444444444444
F1-Score: 0.918918918918919
K = 10
Accuracy: 0.9473684210526315
Precision: 0.8947368421052632
Recall: 0.9444444444444444
F1-Score: 0.918918918918919
K = 11
Accuracy: 0.9649122807017544

Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 12
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 13
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 14
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 15
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 16
Accuracy: 0.9736842105263158
Precision: 0.9459459459459459
Recall: 0.9722222222222222
F1-Score: 0.958904109589041
K = 17
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 18
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 19
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 20
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459


K = 21
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 22
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 23
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 24
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 25
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 26
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 27
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 28
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 29
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473
Recall: 0.9722222222222222
F1-Score: 0.9459459459459459
K = 30
Accuracy: 0.9649122807017544
Precision: 0.9210526315789473

Recall: 0.9722222222222222
F1-Score: 0.9459459459459459

Highest Classification Accuracy Achieved: 97.36842105263158 for K number = 16

```
[30]: # Plot The Graph  
plot_result_with_k(k_list, acc_list, precision_list, recall_list, f1score_list)
```

Plotting classification accuracy, precision, recall and F1-score over a different number of Ks



Problem 5: Gaussian Naive Bayes with PCA – Performance Evaluation

```
[ ]: # Principal Component Analysis (PCA) with Gaussian Naive Bayes (GNB)  
k_list, acc_list, precision_list, recall_list, f1score_list =   
    GaussianNB_model_training_pca(X, Y)
```

K = 1
Accuracy: 0.9122807017543859
Precision: 0.90625
Recall: 0.8055555555555556
F1-Score: 0.8529411764705882
K = 2
Accuracy: 0.8947368421052632
Precision: 0.9
Recall: 0.75
F1-Score: 0.8181818181818182
K = 3
Accuracy: 0.868421052631579
Precision: 0.8387096774193549

Recall: 0.7222222222222222
F1-Score: 0.7761194029850746
K = 4
Accuracy: 0.9035087719298246
Precision: 0.8787878787878788
Recall: 0.8055555555555556
F1-Score: 0.8405797101449275
K = 5
Accuracy: 0.9122807017543859
Precision: 0.8823529411764706
Recall: 0.8333333333333334
F1-Score: 0.8571428571428571
K = 6
Accuracy: 0.9122807017543859
Precision: 0.90625
Recall: 0.8055555555555556
F1-Score: 0.8529411764705882
K = 7
Accuracy: 0.9122807017543859
Precision: 0.90625
Recall: 0.8055555555555556
F1-Score: 0.8529411764705882
K = 8
Accuracy: 0.8947368421052632
Precision: 0.8529411764705882
Recall: 0.8055555555555556
F1-Score: 0.8285714285714286
K = 9
Accuracy: 0.9035087719298246
Precision: 0.8571428571428571
Recall: 0.8333333333333334
F1-Score: 0.8450704225352113
K = 10
Accuracy: 0.8947368421052632
Precision: 0.8333333333333334
Recall: 0.8333333333333334
F1-Score: 0.8333333333333334
K = 11
Accuracy: 0.9035087719298246
Precision: 0.8378378378378378
Recall: 0.8611111111111112
F1-Score: 0.8493150684931506
K = 12
Accuracy: 0.9035087719298246
Precision: 0.8378378378378378
Recall: 0.8611111111111112
F1-Score: 0.8493150684931506
K = 13

Accuracy: 0.9210526315789473
Precision: 0.8648648648648649
Recall: 0.8888888888888888
F1-Score: 0.8767123287671232
K = 14
Accuracy: 0.9122807017543859
Precision: 0.8611111111111112
Recall: 0.8611111111111112
F1-Score: 0.8611111111111112
K = 15
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 16
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 17
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 18
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 19
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 20
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 21
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 22
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112

```

F1-Score: 0.8732394366197183
K = 23
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 24
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 25
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 26
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 27
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 28
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 29
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183
K = 30
Accuracy: 0.9210526315789473
Precision: 0.8857142857142857
Recall: 0.8611111111111112
F1-Score: 0.8732394366197183

```

```

-----
Highest Classification Accuracy Achieved: 92.10526315789474 for K number = 13

```

```

[31]: # Plot The Graph
      plot_result_with_k(k_list, acc_list, precision_list, recall_list, f1score_list)

```

Plotting classification accuracy, precision, recall and F1-score over a different number of Ks

